

Gold Price Prediction

In this notebook we can see how to predict the gold price using the variables like Date, Open, High, Low, Close, Adj Close, Volume, SP_open, SP_high, SP_low, SP_close, SP_Ajclose, SP_volume, DJ_open, DJ_high, DJ_low, DJ_close, DJ_Ajclose, DJ_volume, EG_open, EG_high, EG_low, EG_close, EG_Ajclose, EG_volume, EU_Price, EU_open, EU_high, EU_low, EU_Trend, OF_Price, OF_Open, OF_High, OF_Low, OF_Volume, OF_Trend, OS_Price, OS_Open, OS_High, OS_Low, OS_Trend, SF_Price, SF_Open, SF_High, SF_Low, SF_Volume, SF_Trend, USB_Price, USB_Open, USB_High, USB_Low, USB_Trend, PLT_Price, PLT_Open, PLT_High, PLT_Low, PLT_Trend, PLD_Price, PLD_Open, PLD_High, PLD_Low, PLD_Trend, RHO_PRICE, USDI_Price, USDI_Open, USDI_High, USDI_Low, USDI_Volume, USDI_Trend, GDX_Open, GDX_High, GDX_Low, GDX_Close, GDX_Adj Close, GDX_Volume, USO_Open, USO_High, USO_Low, USO_Close, USO_Adj Close, USO_Volume

Import packages

```
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split , GridSearchCV
from sklearn.ensemble import RandomForestRegressor
from sklearn.decomposition import PCA
from sklearn.metrics import mean_squared_error
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt
import seaborn as sns

/opt/conda/lib/python3.10/site-packages/scipy/__init__.py:146:
UserWarning: A NumPy version >=1.16.5 and <1.23.0 is required for this
version of SciPy (detected version 1.23.5
  warnings.warn(f"A NumPy version >={np_minversion} and
<{np_maxversion}")
```

Import Data set

```
data=pd.read_csv("/kaggle/input/gold-price-prediction-dataset/
FINAL_USO.csv")
df=pd.DataFrame(data)
```

Exploratory Data Analysis

```
df.head()
```

	Date	Open	High	Low	Close	Adj
Close \						
0	2011-12-15	154.740005	154.949997	151.710007	152.330002	
1	2011-12-16	154.309998	155.369995	153.899994	155.229996	
2	2011-12-19	155.479996	155.860001	154.360001	154.869995	
3	2011-12-20	156.820007	157.429993	156.580002	156.979996	
4	2011-12-21	156.979996	157.529999	156.130005	157.160004	

	Volume	SP_open	SP_high	SP_low	...	GDX_Low
GDX_Close \						
0	21521900	123.029999	123.199997	121.989998	...	51.570000
1	18124300	122.230003	122.949997	121.300003	...	52.040001
2	12547200	122.059998	122.320000	120.029999	...	51.029999
3	9136300	122.180000	124.139999	120.370003	...	52.369999
4	11996100	123.930000	124.360001	122.750000	...	52.419998

	GDX_Adj Close	GDX_Volume	USO_Open	USO_High	USO_Low
USO_Close \					
0	48.973877	20605600	36.900002	36.939999	36.049999
1	49.921513	16285400	36.180000	36.500000	35.730000
2	48.490578	15120200	36.389999	36.450001	35.930000
3	50.215282	11644900	37.299999	37.610001	37.220001
4	50.186852	8724300	37.669998	38.240002	37.520000

	USO_Adj Close	USO_Volume
0	36.130001	12616700
1	36.270000	12578800
2	36.200001	7418200
3	37.560001	10041600
4	38.110001	10728000

[5 rows x 81 columns]

df.shape

```
(1718, 81)
```

```
df.describe()
```

	Open	High	Low	Close	Adj Close
\count	1718.000000	1718.000000	1718.000000	1718.000000	1718.000000
mean	127.323434	127.854237	126.777695	127.319482	127.319482
std	17.526993	17.631189	17.396513	17.536269	17.536269
min	100.919998	100.989998	100.230003	100.500000	100.500000
25%	116.220001	116.540001	115.739998	116.052502	116.052502
50%	121.915001	122.325001	121.369999	121.795002	121.795002
75%	128.427494	129.087498	127.840001	128.470001	128.470001
max	173.199997	174.070007	172.919998	173.610001	173.610001

	Volume	SP_open	SP_high	SP_low
SP_close ... \count	1.718000e+03	1718.000000	1718.000000	1718.000000
1718.000000 ...				
mean	8.446327e+06	204.490023	205.372637	203.487014
204.491222 ...				
std	4.920731e+06	43.831928	43.974644	43.618940
43.776999 ...				
min	1.501600e+06	122.059998	122.320000	120.029999
120.290001 ...				
25%	5.412925e+06	170.392498	170.962506	169.577499
170.397500 ...				
50%	7.483900e+06	205.464996	206.459999	204.430000
205.529999 ...				
75%	1.020795e+07	237.292500	237.722500	236.147503
236.889996 ...				
max	9.380420e+07	293.089996	293.940002	291.809998
293.579987 ...				

	GDX_Low	GDX_Close	GDX_Adj Close	GDX_Volume
USO_Open \count	1718.000000	1718.000000	1718.000000	1.718000e+03
1718.000000				
mean	26.384575	26.715012	25.924624	4.356515e+07
22.113417				
std	10.490908	10.603110	9.886570	2.909151e+07
11.431056				
min	12.400000	12.470000	12.269618	4.729000e+06

```

7.820000
25%      20.355000      20.585000      20.180950  2.259968e+07
11.420000
50%      22.870001      23.054999      22.677604  3.730465e+07
16.450000
75%      26.797500      27.317500      26.478154  5.697055e+07
34.419998
max       56.770000      57.470001      54.617039  2.321536e+08
41.599998

```

```

      US0_High      US0_Low      US0_Close  US0_Adj Close
US0_Volume
count  1718.000000  1718.000000  1718.000000    1718.000000
1.718000e+03
mean     22.307148    21.904657    22.109051     22.109051
1.922313e+07
std      11.478671    11.373997    11.432787     11.432787
1.575743e+07
min       8.030000     7.670000     7.960000     7.960000
1.035100e+06
25%      11.500000    11.300000    11.392500     11.392500
6.229500e+06
50%      16.635001    16.040000    16.345000     16.345000
1.613015e+07
75%      34.667499    34.110000    34.417499     34.417499
2.672375e+07
max      42.299999    41.299999    42.009998     42.009998
1.102657e+08

```

[8 rows x 80 columns]

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1718 entries, 0 to 1717
Data columns (total 81 columns):
#   Column              Non-Null Count  Dtype
---  -
0   Date                1718 non-null  object
1   Open                1718 non-null  float64
2   High                1718 non-null  float64
3   Low                 1718 non-null  float64
4   Close               1718 non-null  float64
5   Adj Close           1718 non-null  float64
6   Volume              1718 non-null  int64
7   SP_open              1718 non-null  float64
8   SP_high              1718 non-null  float64
9   SP_low               1718 non-null  float64
10  SP_close             1718 non-null  float64
11  SP_Ajclose           1718 non-null  float64

```

12	SP_volume	1718	non-null	int64
13	DJ_open	1718	non-null	float64
14	DJ_high	1718	non-null	float64
15	DJ_low	1718	non-null	float64
16	DJ_close	1718	non-null	float64
17	DJ_Ajclose	1718	non-null	float64
18	DJ_volume	1718	non-null	int64
19	EG_open	1718	non-null	float64
20	EG_high	1718	non-null	float64
21	EG_low	1718	non-null	float64
22	EG_close	1718	non-null	float64
23	EG_Ajclose	1718	non-null	float64
24	EG_volume	1718	non-null	int64
25	EU_Price	1718	non-null	float64
26	EU_open	1718	non-null	float64
27	EU_high	1718	non-null	float64
28	EU_low	1718	non-null	float64
29	EU_Trend	1718	non-null	int64
30	OF_Price	1718	non-null	float64
31	OF_Open	1718	non-null	float64
32	OF_High	1718	non-null	float64
33	OF_Low	1718	non-null	float64
34	OF_Volume	1718	non-null	int64
35	OF_Trend	1718	non-null	int64
36	OS_Price	1718	non-null	float64
37	OS_Open	1718	non-null	float64
38	OS_High	1718	non-null	float64
39	OS_Low	1718	non-null	float64
40	OS_Trend	1718	non-null	int64
41	SF_Price	1718	non-null	int64
42	SF_Open	1718	non-null	int64
43	SF_High	1718	non-null	int64
44	SF_Low	1718	non-null	int64
45	SF_Volume	1718	non-null	int64
46	SF_Trend	1718	non-null	int64
47	USB_Price	1718	non-null	float64
48	USB_Open	1718	non-null	float64
49	USB_High	1718	non-null	float64
50	USB_Low	1718	non-null	float64
51	USB_Trend	1718	non-null	int64
52	PLT_Price	1718	non-null	float64
53	PLT_Open	1718	non-null	float64
54	PLT_High	1718	non-null	float64
55	PLT_Low	1718	non-null	float64
56	PLT_Trend	1718	non-null	int64
57	PLD_Price	1718	non-null	float64
58	PLD_Open	1718	non-null	float64
59	PLD_High	1718	non-null	float64
60	PLD_Low	1718	non-null	float64

```

61 PLD_Trend      1718 non-null int64
62 RHO_PRICE     1718 non-null int64
63 USDI_Price    1718 non-null float64
64 USDI_Open     1718 non-null float64
65 USDI_High     1718 non-null float64
66 USDI_Low      1718 non-null float64
67 USDI_Volume   1718 non-null int64
68 USDI_Trend    1718 non-null int64
69 GDX_Open      1718 non-null float64
70 GDX_High      1718 non-null float64
71 GDX_Low       1718 non-null float64
72 GDX_Close     1718 non-null float64
73 GDX_Adj Close 1718 non-null float64
74 GDX_Volume    1718 non-null int64
75 US0_Open      1718 non-null float64
76 US0_High      1718 non-null float64
77 US0_Low       1718 non-null float64
78 US0_Close     1718 non-null float64
79 US0_Adj Close 1718 non-null float64
80 US0_Volume    1718 non-null int64
dtypes: float64(58), int64(22), object(1)
memory usage: 1.1+ MB

```

```
df.nunique()
```

```

Date      1718
Open      1363
High      1368
Low       1356
Close     1384
...
US0_High  1102
US0_Low   1124
US0_Close 1128
US0_Adj Close 1128
US0_Volume 1716
Length: 81, dtype: int64

```

```

X = data.drop(['Date', 'Close'], axis=1) # Remove 'Date' and 'Close'
columns
y = data['Close']

```


Analysis by Charts

```

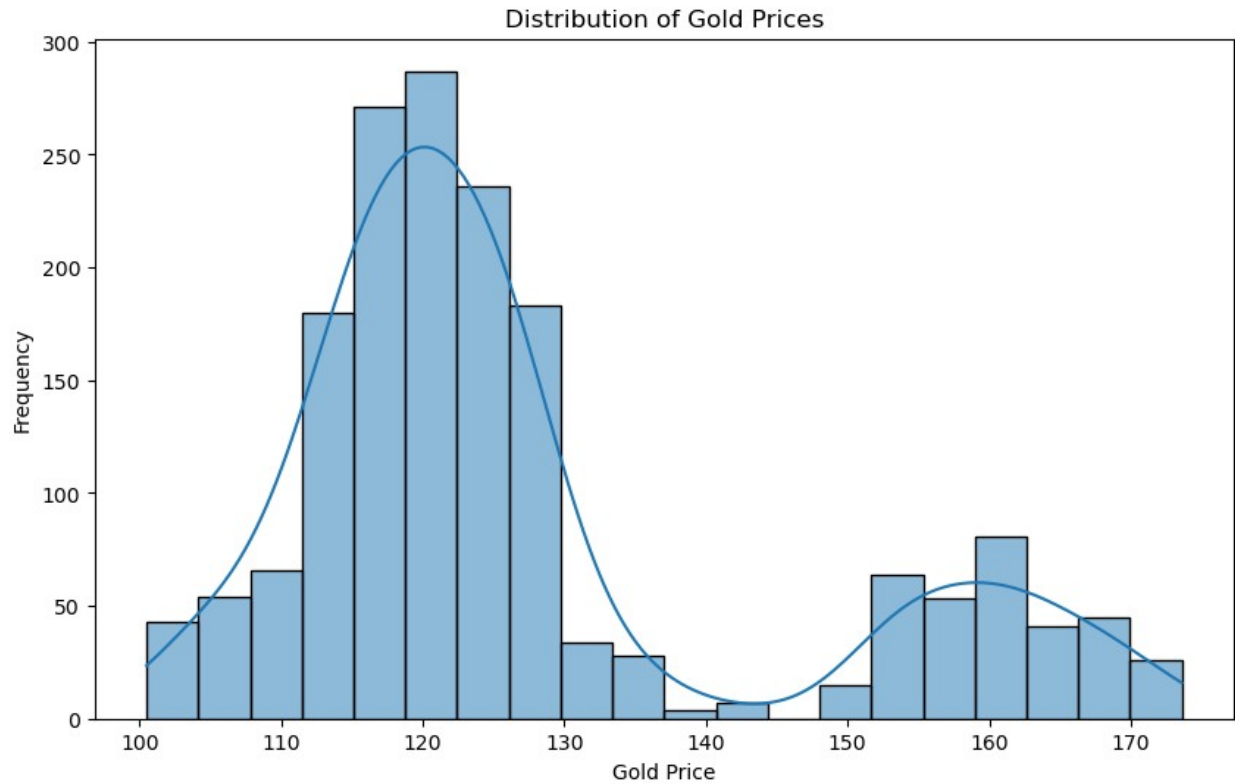
# Line Chart - Gold Price over Time
plt.figure(figsize=(12, 6))
plt.plot(data['Date'], data['Close'], marker='o')
plt.xlabel('Date')

```

```
plt.ylabel('Gold Price')
plt.title('Gold Price Over Time')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



```
# Histogram - Distribution of Gold Prices
plt.figure(figsize=(10, 6))
sns.histplot(data['Close'], bins=20, kde=True)
plt.xlabel('Gold Price')
plt.ylabel('Frequency')
plt.title('Distribution of Gold Prices')
plt.show()
```

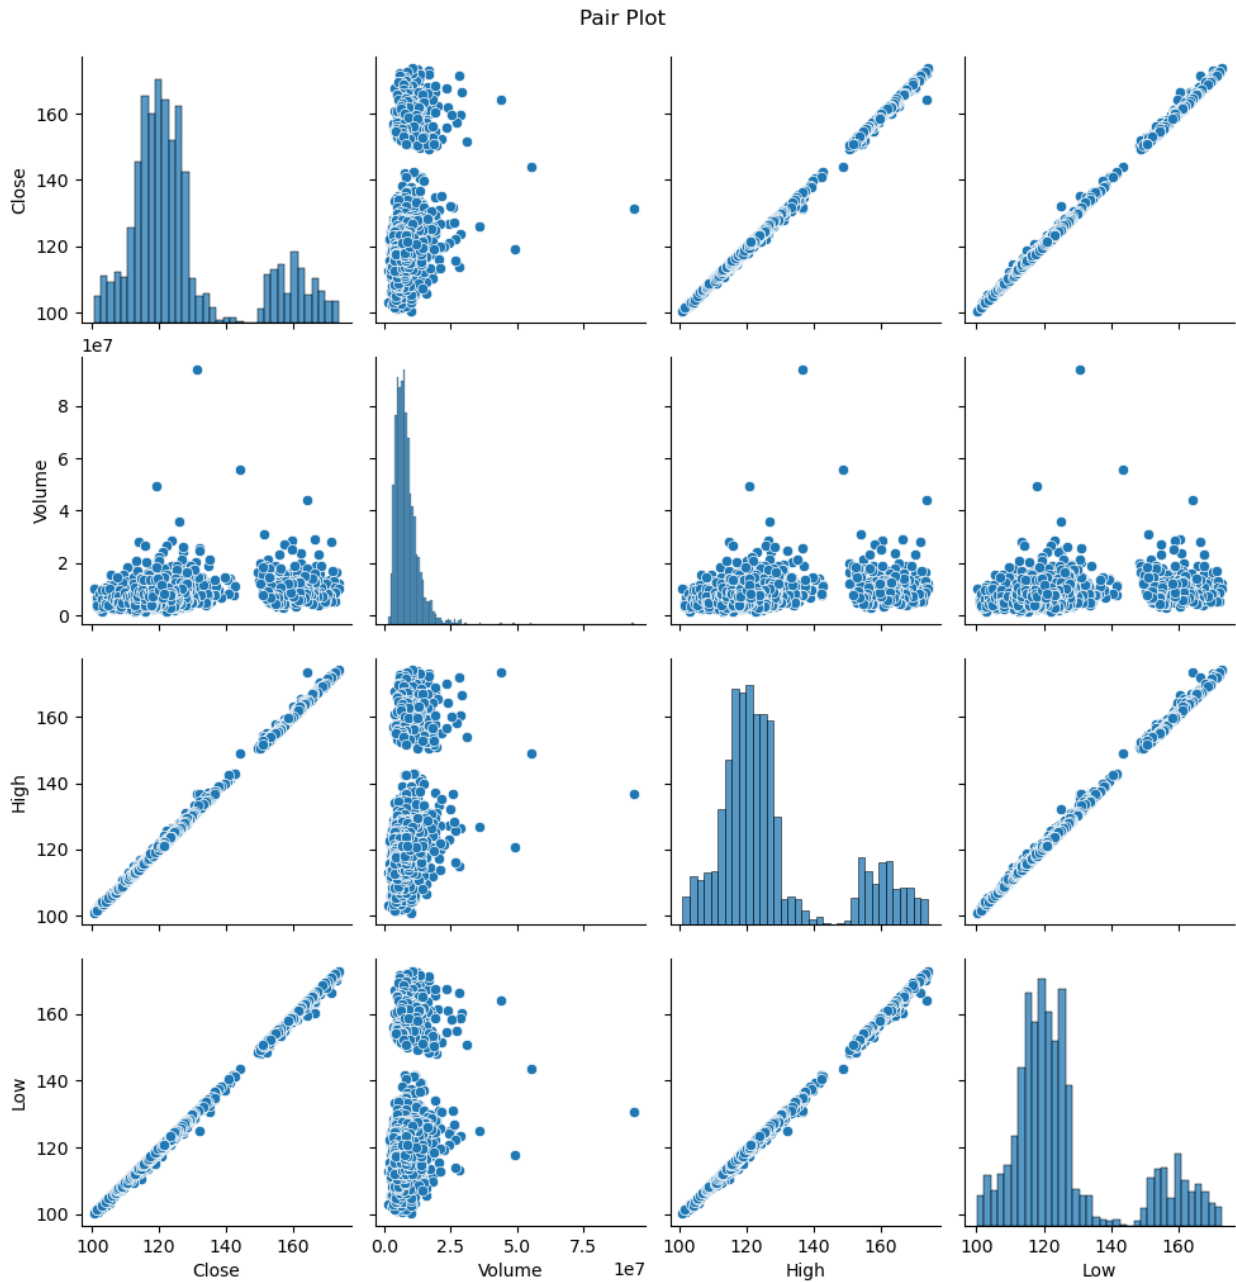


```
# Scatter Plot - Gold Price vs. Volume
plt.figure(figsize=(10, 6))
plt.scatter(data['Volume'], data['Close'], alpha=0.5)
plt.xlabel('Volume')
plt.ylabel('Gold Price')
plt.title('Gold Price vs. Volume')
plt.show()
```





```
# Pair Plot
sns.pairplot(data[['Close', 'Volume', 'High', 'Low']])
plt.suptitle('Pair Plot', y=1.02)
plt.show()
```

/opt/conda/lib/python3.10/site-packages/seaborn/axisgrid.py:118:
UserWarning: The figure layout has changed to tight
self._figure.tight_layout(*args, **kwargs)



```
# Box Plot - Gold Price by Trend
plt.figure(figsize=(10, 6))
sns.boxplot(x=data['PLD_Trend'], y=data['Close'])
plt.xlabel('Trend')
plt.ylabel('Gold Price')
plt.title('Gold Price by Trend')
plt.show()
```




Correlation between the variables

```
# Select only the numeric columns for correlation calculation
numeric_columns = data.select_dtypes(include=[np.number])

# Calculate the correlation matrix
corr_matrix = numeric_columns.corr()
```

HeatMap

```
# Create a heatmap using seaborn
plt.figure(figsize=(10, 8))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt=".2f",
center=0)
plt.title('Correlation Heatmap')
plt.show()
```



Split the data into training and testing sets

```
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=0)
```

Define Preprocessing Step

```
# Scale features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Apply PCA for dimensionality reduction
n_components = 30 # Choose an appropriate number of components
```

```
pca = PCA(n_components=n_components)
X_train_pca = pca.fit_transform(X_train_scaled)
X_test_pca = pca.transform(X_test_scaled)
```

Creating and Training Model

```
param_grid = {
    'n_estimators': [100, 200],
    'max_depth': [None, 10],
    'min_samples_split': [2, 5],
    'min_samples_leaf': [1, 2]
}

grid_search = GridSearchCV(RandomForestRegressor(random_state=0,
n_jobs=-1), param_grid, cv=5)
grid_search.fit(X_train_pca, y_train)
best_model = grid_search.best_estimator_

/opt/conda/lib/python3.10/site-packages/scipy/__init__.py:146:
UserWarning: A NumPy version >=1.16.5 and <1.23.0 is required for this
version of SciPy (detected version 1.23.5
  warnings.warn(f"A NumPy version >={np_minversion} and
<{np_maxversion}")
/opt/conda/lib/python3.10/site-packages/scipy/__init__.py:146:
UserWarning: A NumPy version >=1.16.5 and <1.23.0 is required for this
version of SciPy (detected version 1.23.5
  warnings.warn(f"A NumPy version >={np_minversion} and
<{np_maxversion}")
/opt/conda/lib/python3.10/site-packages/scipy/__init__.py:146:
UserWarning: A NumPy version >=1.16.5 and <1.23.0 is required for this
version of SciPy (detected version 1.23.5
  warnings.warn(f"A NumPy version >={np_minversion} and
<{np_maxversion}")
/opt/conda/lib/python3.10/site-packages/scipy/__init__.py:146:
UserWarning: A NumPy version >=1.16.5 and <1.23.0 is required for this
version of SciPy (detected version 1.23.5
  warnings.warn(f"A NumPy version >={np_minversion} and
<{np_maxversion}")
```

Make predictions on the test data

```
y_pred = best_model.predict(X_test_pca)

# Calculate Mean Squared Error
mse = mean_squared_error(y_test, y_pred)
print(f"Mean Squared Error: {mse}")
```


Mean Squared Error: 1.956613975836765

Plot the actual vs. predicted values

```
plt.figure(figsize=(10, 6))
plt.scatter(y_test.index, y_test, color='blue', label='Actual')
plt.scatter(y_test.index, y_pred, color='red', label='Predicted')
plt.xlabel('Index')
plt.ylabel('Gold Price')
plt.title('Gold Price Prediction')
plt.legend()
plt.show()
```

